

Host a Website

on AWS with Terraform (Infrastructure as Code)



Lucas Nebelung
42304499

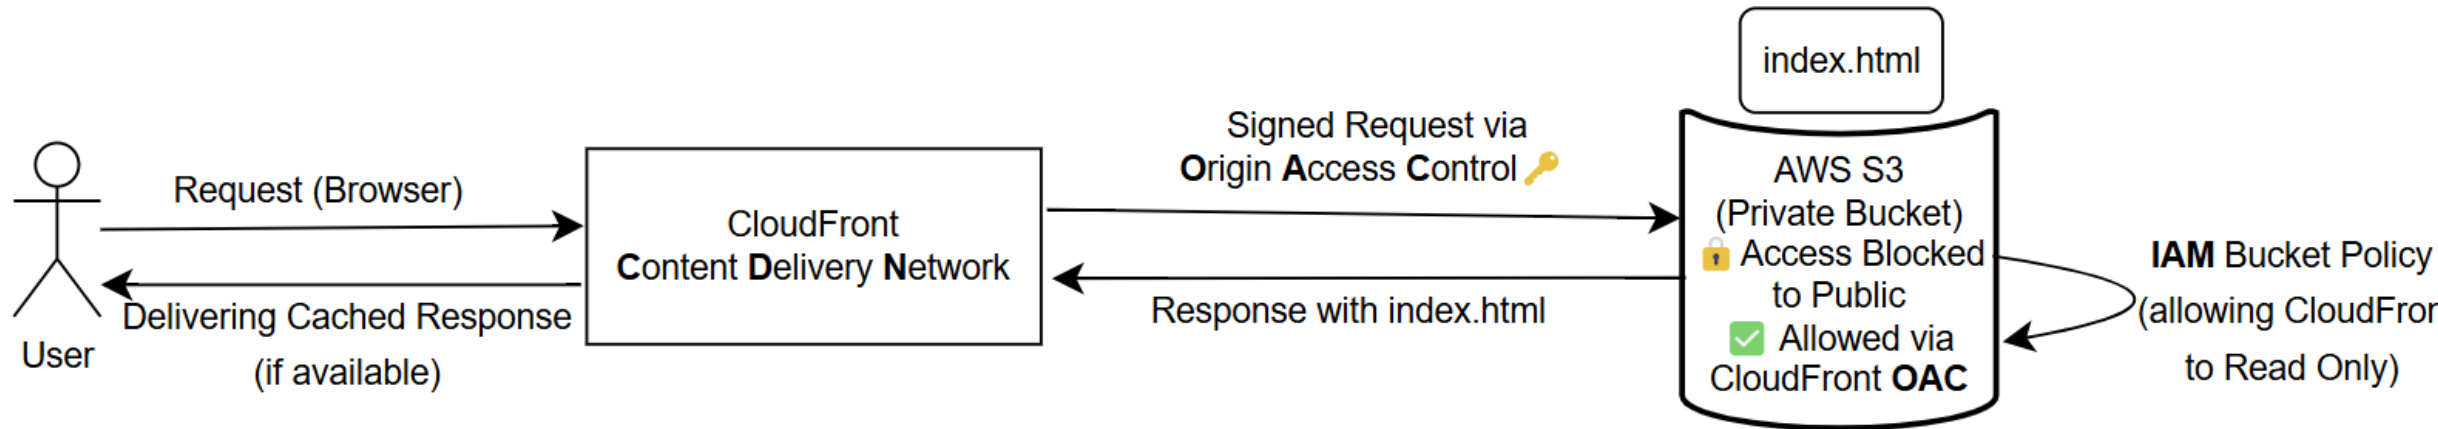
Overview

- **Goal:** Host a static “Hello World” webpage on *AWS* using *IaC*
- **Requirements:**
 - Global availability
 - Low latency
 - Autoscaling

(optional) secure, cost-efficient, modular



Solution: Architecture



Why this Architecture?



CloudFront + S3 REST API + OAC (Private Bucket)

- More secure
- BUT: More complex
-> requires OAC, bucket policy and CloudFront Config

CloudFront + S3 Website Hosting (Public Bucket)

- Less secure
-> Bucket must be public for CloudFront to fetch content
- Simpler

Why S3 + CloudFront?

- **S3**: Simple, cost-effective storage for static content.
- **CloudFront**: Provides a global edge network (CDN) for low-latency delivery.
- **Private Bucket + OAC**: Ensures the bucket is *not* publicly accessible; only CloudFront can read objects.
- **Autoscaling**
S3 and CloudFront automatically scale based on traffic load. No manual configuration needed for extra servers.

Why Terraform?

- **Infrastructure as Code (IaC):** All resources defined in *.tf* files
- **Version Control:** Easily track changes and collaborate.
- **Consistency:** The same configuration can be applied across environments
- **Cloud-Agnostic:** Could adapt to other clouds in the future if needed.

Code Explanation I: S3 Bucket

- 1) selecting AWS region to be Frankfurt (Germany)
- 2)
 - Creates a private S3 bucket named *iu-test-bucket* (No ACL or website hosting block so the bucket defaults to private)
 - Uploads local *index.html* into the newly created private S3 bucket.

```
#####  
##### Hosting a simple webpage on AWS using Terraform #####  
#####  
  
#####  
# 1) Terraform Configuration & Provider Settings  
#####  
  
provider "aws" {  
  region = "eu-central-1"  
}  
  
#####  
# 2) create a private S3 Bucket  
#####  
  
# create an S3 bucket with name  
resource "aws_s3_bucket" "private_bucket" {  
  bucket = "iu-test-bucket"  
}  
  
# upload the index.html file to the S3 bucket  
resource "aws_s3_object" "index" {  
  bucket = aws_s3_bucket.private_bucket.id  
  key    = "index.html"      # name of the file in S3  
  source = "index.html"      # local file  
  content_type = "text/html"  
}
```

Code Explanation II: OAC

- Defining **Origin Access Control (OAC)** for the CloudFront Distribution
- Setting “*always*” and “*sigv4*” for maximum security

```
#####  
# 3) Origin Access Control (OAC)  
#####  
  
resource "aws_cloudfront_origin_access_control" "oac" {  
    name                        = "my-s3-oac"  
    origin_access_control_origin_type = "s3"  
    signing_behavior            = "always"  
    signing_protocol            = "sigv4"  
}
```


Code Explanation III: CloudFront Distribution

- Creating **CloudFront Distribution** and setting *index.html* as default
- Specifying origin to our "*private-s3-origin*" and assigning previously specified **OAC**
- Setting *default_cache_behavior* to ["GET", "HEAD"] as sufficient for our use case
- Forcing all traffic to use HTTPS, improving security
- Using simple settings for *forwarded_values*, *restrictions* and *viewer_certificate* (so website available via *.cloudfront.net)

```
#####
# 4) Cloudfront Distribution with S3 Bucket as Origin
#####

#create a cloudfront distribution first
resource "aws_cloudfront_distribution" "cdn" {
  enabled = true
  default_root_object = "index.html"
  is_ipv6_enabled = true
  wait_for_deployment = true

  # Where does Cloudfront pull content from?
  origin {
    domain_name = aws_s3_bucket.private_bucket.bucket_regional_domain_name
    origin_id = "private-s3-origin"
    origin_access_control_id = aws_cloudfront_origin_access_control.oac.id
  }

  #set cache behaviors
  default_cache_behavior {
    target_origin_id = "private-s3-origin"
    viewer_protocol_policy = "redirect-to-https"
    allowed_methods = ["GET", "HEAD", ]
    cached_methods = ["GET", "HEAD",]
    compress = true

    # simple default setting for query string, cookies, and headers
    forwarded_values {
      query_string = false
      cookies {
        forward = "none"
      }
    }
  }

  #no geo or specific HTTPS, SSL/TLS encryption restrictions
  restrictions {
    geo_restriction {
      restriction_type = "none"
    }
  }

  viewer_certificate {
    cloudfront_default_certificate = true
  }
}
```

Code Explanation IV: IAM Bucket Policy

- Fetching current AWS account
- Creating **IAM** policy document
 - > allows CloudFront to READ (GET) only
 - > allows CloudFront to access all objects inside specified bucket
- *Principals* ensures only CloudFront (and no other Service) can access
- *Condition* enforces requests come from our AWS Account
- Assigning created Bucket Policy to our Bucket

```
#####  
# 5) Bucket Policy to Allow CloudFront Read  
#####  
  
# get the current AWS account identity and assign it to "current"  
data "aws_caller_identity" "current" {}  
  
# create a IAM policy document  
data "aws_iam_policy_document" "bucket_policy_doc" {  
  statement {  
    sid = "AllowCloudFrontServicePrincipal" #Statement ID  
    effect = "Allow"  
    actions = ["s3:GetObject"] #Cloudfront can on  
    resources = [  
      "${aws_s3_bucket.private_bucket.arn}/*" #allows Cloud  
    ]  
  
    # only allow CloudFront to access the S3 private bucket  
    principals {  
      type = "Service"  
      identifiers = ["cloudfront.amazonaws.com"]  
    }  
  
    # only allow access from the same AWS account  
    condition {  
      test = "StringEquals"  
      variable = "AWS:SourceAccount"  
      values = [data.aws_caller_identity.current.account_id]  
    }  
  }  
}  
  
# apply the policy to the bucket  
resource "aws_s3_bucket_policy" "bucket_policy" {  
  bucket = aws_s3_bucket.private_bucket.id #specifies which  
  policy = data.aws_iam_policy_document.bucket_policy_doc.json  
}
```

Code Explanation V: Output

- Ensuring Terraform states the *domain* of Website and corresponding bucket after deployment

```
#####  
# 6) Outputs  
#####  
  
# Output the CloudFront domain name  
output "cloudfront_domain_name" {  
    description = "The Website can be accesed with this URL"  
    value       = aws_cloudfront_distribution.cdn.domain_name  
}  
  
# Output the S3 bucket name  
output "bucket_name" {  
    value = aws_s3_bucket.private_bucket.bucket  
}
```

Frameworks used ...

- **AWS** Documentation
<https://docs.aws.amazon.com/>
-> link to all relevant services including S3, CloudFront, IAM
- **Terraform** Documentation
<https://developer.hashicorp.com/terraform/docs>
-> link to Terraform Installation, commands etc